

CS392 Database System Concept

Assignment 7

Due April 28th, 2014

1. (15') Given key values: (2, 3, 5, 7, 11, 19, 23, 29, 31). Construct B⁺-Tree under the following three cases (assume the key is inserted in ascending order)
 - a. Four pointer per node
 - b. Six pointer per node
 - c. Eight pointer per node
2. (15') For each case in problem 1, show the tree after each of the following series of operations:
 - a. Insert 9
 - b. Insert 10
 - c. Insert 8
 - d. Delete 23
 - e. Delete 19
3. (15') Suppose there is a relation R(A, B, C), with B⁺ tree index with search key (A, B).
 - a. What is the worst case cost of finding records satisfying $10 < A < 50$ using this index, in terms of the number of records retrieved n_1 and the height h of the tree?
 - b. What is the worst case cost of finding records satisfying $10 < A < 50 \wedge 5 < B < 10$ using this index, in terms of the number of records n_2 that satisfy this selection, as well as n_1 and h defined above.
 - c. Under what condition on n_1 and n_2 would the index be an efficient way of finding records satisfying $10 < A < 50 \wedge 5 < B < 10$.
4. (20') Let relations $r_1(A, B, C)$ and $r_2(C, D, E)$ have the following properties: r_1 has 50,000 tuples, r_2 has 45,000 tuples, 25 tuples of r_1 fit on one block, and 30 tuples of r_2 fit on one block. Estimate the number of block transfers and seeks required, using each of the following join strategies for $r_1 \bowtie r_2$
 - a. Nested-loop join
 - b. Block nested-loop join
 - c. Merge join
 - d. Hash join
5. (15') Suppose that a B⁺-tree index on building is available on relation department, and that no other index is available. What would be the best way to handle the following selections that involve negation?
 - a. $\sigma_{\neg(\text{building} < \text{"Watson"})}(\text{department})$
 - b. $\sigma_{\neg(\text{building} = \text{"Watson"})}(\text{department})$
 - c. $\sigma_{\neg(\text{building} < \text{"Watson"} \vee \text{budget} < 50000)}(\text{department})$

6. (20') Pipelining is used to avoid writing intermediate results to disk. Suppose you need to sort relation r using sort-merge and merge-join the result with an already sorted relations.
- Describe how the output of the sort of r can be pipelined to the merge join without being written back to disk.
 - The same idea is applicable even if both inputs to the merge-join are the outputs of sort-merge operations. However, the available memory has to be shared between the two merge operations (the merge-join algorithm itself needs very little memory). What is the effect of having to share memory on the cost of each sort-merge operation.