

## Assignment 8

Qinglin Li, 5110309074

### Problem 1

Use the B+-tree to find the first tuple with *branch\_name* = "Downtown" and *branch\_city* < "Brooklyn", then scan the following tuples and check whether each tuple satisfies the condition on *assets*. Stop the scanning when *branch\_name* ≠ "Downtown" and *branch\_city* ≥ "Brooklyn"

### Problem 2

In a **serial schedule**, no transaction starts until a running transaction has ended.  
A **serializable schedule** is a schedule that is equivalent (in its outcome) to a serial schedule

### Problem 3

- a. If we execute T1 first, the result should be  $A = 0$  and  $B = 1$   
If we execute T2 first, the result should be  $A = 1$  and  $B = 0$   
So in both case, the consistency requirement is satisfied.

b.

T1	T2
Read(A)	
	Read(B)
Read(B)	
	Read(A)
If A=0 then B:=B+1	
	If B=0 then A:=A+1
Write(B)	
	Write(A)

c. No.

Because If we want to produce a serializable result, we must execute the first read instruction after the other's write.

### Problem 4

1. To improve throughput and resource utilization.
2. To Reduce waiting time.

## Problem 5

a. T31:

```
lock-S(A)
read(A)
lock-X(B)
read(B)
if A = 0 then B:= B + 1
write(B)
unlock(A)
unlock(B)
```

T32:

```
lock-S(B)
read(B)
lock-X(A)
read(A)
if B = 0 then A := A + 1
write(A)
unlock(B)
unlock(A)
```

b. Yes.

T31 acquire the lock on A first and then T32 acquire the lock on B.

## Problem 6

Advantage: It produces only cascadeless schedules, recovery is very easy.

Disadvantage: The set of schedules obtainable is a subset of those obtainable from plain two phase locking, thus concurrency is reduced

## Problem 7

Suppose a transaction is rolled back because of a newer transaction's reading or writing the data which it plans to write. If the rolled back transaction is re-introduced with the same timestamp, the transaction should be rolled back again because of the same reason and the process would never end.

## Problem 8

A transaction may become the victim of deadlock-prevention rollback arbitrarily many times, thus creating a potential starvation situation.

## Problem 9

Volatile storage is storage which fails after a power failure. Caches, main memories are volatile storage.

Non-volatile storage is storage which retains its content after power failures. Magnetic disk, tapes are non-volatile storage.

Stable storage is storage which survives any kind of failure. This type of storage can only be approximated by replicating data.

Volatile memory is the fastest and non-volatile storage is slower. Stable storage is the slowest because of the data replication.

## Problem 10

- a. Stable storage cannot really be implemented because all storage devices are made of hardware, and all hardware is vulnerable to mechanical or electronic device failures.
- b. Database systems approximate stable storage by writing data to multiple storage devices simultaneously. Even if one of the devices crashes, the data will still be available on a different device.

## Problem 11

In a system crash, the CPU goes down, and disk may also crash. But stable-storage at the site is assumed to survive system crashes.

In a disaster, everything at a site is destroyed.